

Shiksha Setu: A Measurement-Driven, Local-First AI Tutoring Platform for Multilingual Indian Education

Rachit Ranka, Aaditya Saini, Priyanka Devi, Yogesh Kohli, Ruhani Mehta
Department of Computer Science Engineering
Chandigarh University, Mohali 140413, Punjab, India
rankarachit5@gmail.com, aadityasaini899@gmail.com, pjangra9105@gmail.com,
ykohli2005@gmail.com, ruhanimehta019@gmail.com

Abstract—Existing AI tutoring systems depend on cloud infrastructure, operate primarily in English, and are not aligned to a national curriculum. This paper presents Shiksha Setu, a local-first tutoring platform for the Indian NCERT curriculum in which the corpus, its embeddings and the whole of retrieval are local. Generation is a hosted call, and the paper is explicit about which half is which, and does not claim to be offline. All results are measured against 153 ingested textbooks spanning classes 1–12 (42,995 indexed passages). Retrieval is evaluated by a five-configuration ablation over 22 queries in eleven Indian languages against an English-medium corpus: a lexical baseline retrieves *nothing at all*, because a Devanagari or Perso-Arabic query shares no tokens with an English index. Dense retrieval alone reaches 11 of 22, query rewriting alone 14, and pooling both, the deployed configuration, at 16 ± 1 . A grounding score separates answers written from their retrieved passages from answers that drifted by 0.360 in cosine similarity over 256 pairs, at a threshold whose previously shipped value admitted 35 of 240 negatives. Half-precision inference halves the embedding footprint to 1,083 MB at a mean cosine fidelity of 0.999998, verified by kernel-level peak RSS, enabling deployment on 4 GB machines. Ingestion detects two classes of corpus corruption, including legacy Devanagari fonts that render 66% of Hindi textbooks as Latin gibberish.

Index Terms—Retrieval-Augmented Generation, Multilingual Education, Local-First AI, Indian Languages, Curriculum-Aligned Tutoring, Accessibility, NCERT

I. INTRODUCTION

India is home to 1.4 billion people speaking 22 constitutionally recognized languages [22], written in more than a dozen scripts. The National Council of Educational Research and Training (NCERT) publishes the curriculum that governs instruction for the 24.7 crore children enrolled in school [21], yet its textbooks exist only in English, Hindi, and Urdu. A Tamil-speaking student in rural Tamil Nadu, a Marathi-speaking student in Maharashtra, and a Bengali-speaking student in West Bengal all study the same curriculum, but none of them can read it in their mother tongue unless a human teacher translates it in real time.

Artificial intelligence can bridge this gap, but the dominant paradigm of cloud-hosted large language models introduces three structural failures for Indian education:

Language exclusion. Most educational AI systems operate in English. English is a first language for a negligible share of

the population and a second or third language for roughly one Indian in ten [22], so for the great majority such systems are inaccessible at the point of need. **Infrastructure dependence.** Cloud solutions need consistent bandwidth and bill per query, which makes sustained use impractical in the towns and villages where most students are. **Pedagogical misalignment.** A general-purpose model has no notion of curriculum structure or grade-appropriate complexity, and a class 6 student asking about chemical reactions needs a different answer from a class 12 student on the same topic.

This paper presents Shiksha Setu, a platform that addresses the first and third failures and *partially* the second, through a measurement-driven engineering approach.

The qualification is deliberate. Retrieval is entirely local. The corpus, its embeddings and the vector index never leave the machine, but generation is served by a hosted endpoint, so a question’s text does leave it and does incur a per-query cost. The system therefore removes the dependency for the component that dominates storage and latency without removing it altogether. Local generation is configured and available (Section III), but a language model cannot share a 4 GB machine with the embedder, so the offline configuration and the 4 GB configuration are alternatives, not the same system. We state this plainly because “local-first” is otherwise read as “offline”, and the measurements in this paper were taken in the hosted configuration. Every claim is backed by an experiment whose procedure, data, and result are reported, including experiments that produced wrong results before being corrected. The contributions are:

This work contributes an ingestion pipeline covering all 559 NCERT textbooks, which detects and rejects the legacy-font corruption that renders 27 of 41 Hindi books unreadable, a defect found only by measuring for it. It contributes a cross-lingual retrieval system that answers questions in eleven Indian languages from an English corpus, with a controlled ablation over 22 queries establishing that query rewriting and direct embedding fail on disjoint inputs, so that pooling them retrieves correctly for 16 ± 1 of 22 where either alone manages 11 or 14 and a lexical baseline manages none. It contributes a memory strategy verified by kernel-level measurement, which

brings the serving pipeline inside a 4 GB budget and corrects an earlier arithmetic estimate that understated consumption by 31%. And it contributes a grounding metric that detects when fluent, confident answers have drifted from their source material, a failure indistinguishable from a correct answer without quantitative measurement, calibrated here on 256 answer and passage pairs instead of hand-checked examples and compared against three alternative estimators.

II. RELATED WORK

Lewis et al. [2] introduced retrieval-augmented generation (RAG) for knowledge-intensive NLP tasks, since surveyed at length [19], establishing the paradigm of conditioning language model output on retrieved passages. Shiksha Setu adopts this architecture with a critical modification: a grade-level filter inserted between retrieval and generation, because the same topic at class 6 and class 12 requires different explanations. Karpukhin et al. [3] demonstrated the superiority of dense passage retrieval over BM25. Ma et al. [4] showed that query rewriting before embedding improves retrieval quality. The technique that resolves the romanized Hindi failure. Nogueira and Cho [5] established passage re-ranking with cross-encoder models. The BGE-Reranker component exists in the codebase but is not yet wired into the retrieval path.

Chen et al. [6] presented BGE-M3, a multilingual embedding model supporting over 100 languages in a unified 1024-dimensional vector space. Its cross-lingual property, that semantically equivalent passages in different languages map to nearby vectors, allows a Hindi question to retrieve an English passage without explicit translation. Gala et al. [7] developed IndicTrans2 covering all 22 scheduled Indian languages, serving as the translation backbone. Kakwani et al. [1] provided monolingual corpora and evaluation benchmarks for Indian languages. Khanuja et al. [8] trained MuRIL on transliterated text, directly relevant to the romanized Hindi failure.

The retrieval layer uses the HNSW graph structure proposed by Malkov and Yashunin [11], implemented through the pgvector extension for PostgreSQL. The index parameters ($m = 16$, $ef_construction = 64$) balance build time against recall. The resulting search latency is reported in Fig. 1, not here. Reimers and Gurevych [12] established the sentence-level embedding formulation used to load BGE-M3.

Micikevicius et al. [13] established the safety of half-precision arithmetic for neural network inference, providing the basis for the fp16 optimization that halves the BGE-M3 memory footprint.

Table I compares this work with the systems an Indian student is most likely to encounter: Khanmigo [14], Khan Academy’s GPT-4 tutor; BYJU’S, the largest Indian ed-tech platform; and Bhashini [15], the Government of India’s national language-technology mission.

The comparison is confined to *architectural properties*. Every cell is a structural fact, a figure the vendor has published, or *n/d* where they have not, and no row sets a measurement of ours against a number we could not obtain. Commercial

products publish neither latency distributions nor retrieval quality nor memory footprints, and a table of plausible estimates for them would be a table of invented numbers set in the typography of measurement.

TABLE I
COMPARISON WITH DEPLOYED SYSTEMS. *n/d*: NOT DISCLOSED BY THE OPERATOR.
ROWS ARE ARCHITECTURAL PROPERTIES, NOT PERFORMANCE CLAIMS.

Property	Khanmigo	BYJU’S	Bhashini	Shiksha Setu
NCERT-aligned corpus	no	yes	—	yes
Indian languages	no	n/d	22	11
Corpus held locally	no	no	—	yes
Retrieval runs locally	no	no	—	yes
Generation runs locally	no	no	—	no
Runs without a network	no	no	no	no
Brahmic readability aid	no	no	no	yes
Per-answer grounding	n/d	n/d	—	yes
Corpus corruption audit	n/d	n/d	—	yes
Stated memory target	n/d	n/d	n/d	4 GB
Tutoring function	yes	yes	no	yes

Three rows are stated against our own interest: generation is not local, the system needs a network, and BYJU’S covers our curriculum at a scale we do not approach. What distinguishes this work is the rest. The corpus, index and retrieval are local, a grounding score accompanies every answer, and the pipeline audits its own sources instead of embedding whatever it extracts. We have found no comparable system reporting the last two, which is a claim about the published record and not about what those systems do internally.

III. SYSTEM ARCHITECTURE

A student’s question traverses an eight-stage pipeline from input to response, six of them mandatory. The latencies below were taken on the development machine; what the platform *requires* is the 2,094 MB of Section VI, which is what puts it inside 4 GB. Fig. 1 shows those six with their measured latencies on an 8 GB Apple M1, arranged by the property this paper is about, which is which side of the machine boundary each one runs on.




Fig. 1. The mandatory request path. Medians over 15 warm requests, five questions run three times on an 8 GB Apple M1. The dashed line is the machine boundary: the corpus and all retrieval stay left of it. Generation dominates and is least predictable, 4.2 s median against a 95th percentile of 13.2 s, because it is hosted. All of retrieval is 0.11 s. The grounding check carries no latency because it was not measured separately.

The platform integrates twenty-eight model identifiers, of which **five** are exercised on the critical path and measured in this work. The remainder are configured alternatives that this evaluation does not characterize. The five are BGE-M3 (568 M parameters, embeddings and grounding, the only one resident locally), and four hosted models: LLaMA-3.1-8B for generation and query rewriting, FLUX.1-dev for illustration, Nemotron-Nano-12B for vision OCR, and LLaMA-3.3-70B, compared and not adopted.

BGE-M3 is used twice for distinct purposes: retrieving passages, and later scoring how far an answer strayed from them (Section V-D). Both uses depend on the same property, cross-lingual embedding in a shared space, which is what permits a Devanagari passage and an English answer to be compared numerically at all.

The choice of the 8B model over the 70B variant was driven by measurement: on one fixed prompt at approximately 70 completion tokens against the same endpoint, LLaMA-3.1-8B returned in 0.9 s where LLaMA-3.3-70B took 17–125 s, highly variable.

The persistence layer combines PostgreSQL 17 with pgvector for HNSW-indexed similarity search. Response caching is tiered, and cache keys incorporate a SHA-256 hash of the authorization header so that one student’s answer cannot be served to another.

An earlier migration stored embeddings as `double precision[]` instead of `vector(1024)`, making index creation impossible and condemning every search to a sequential scan.

The 4 GB target is a claim about the machines schools own, and those are predominantly Windows desktops, not the Apple hardware this work was developed on. Portability is therefore part of the claim, not an afterthought.

Device selection is resolved at load time in the order CUDA, Metal (MPS), CPU, and half precision is selected per device (Section VI). Everything the platform needs to know about its host (installed memory, available memory, swap, peak resident set) is read through a single abstraction over `psutil`, which reports identically on Linux, Windows and macOS.

All measurements in this paper were taken on an Apple M1 with 8 GB of unified memory. We report no figures for hardware we did not run on.

IV. CORPUS CONSTRUCTION

NCERT encodes every textbook with a five-character code: `jesc1` represents Class 10, English medium, Science, book 1. The subject letters are not algorithmically derivable. Science is `sc` for classes 7–10 but `cu` (Curiosity) for class 6, and splits into `ph`, `ch`, `bo` at classes 11–12. The catalog is scraped from NCERT’s textbook picker and cached locally.

The complete catalog contains **559 textbooks**: 209 in English, 191 in Hindi, and 159 in Urdu. The platform teaches from a deliberate subset of **263 books**, every English edition, plus the 54 subjects that exist only in Devanagari.

The pipeline processes each book through five stages: ZIP download, per-chapter PDF extraction via PyMuPDF, text

cleanup, semantic chunking (1200 characters, 200-character overlap, paragraph-boundary preference), BGE-M3 embedding (fp16), and storage in PostgreSQL. Each book is committed as a single database transaction for resumability. Download and embedding contend for no shared resource. One waits on a remote web server, the other saturates the GPU. Executed sequentially they idle in turn, so the next book’s archive is fetched on a prefetch thread while the current book is embedded.

The speedup was established by a controlled comparison, not inferred. The same three books (classes 6–8 *Curiosity*) were ingested twice, once with prefetching disabled through `INGEST_PREFETCH=0`, producing byte-identical output in both runs (37 chapters, 1,130 chunks): **193 s per book sequentially against 140 s per book with prefetching, a 1.38× speedup**. Extrapolated, the 153 obtainable books require approximately 6.0 hours on a single Apple M1. The figure is quoted over the obtainable books instead of the full 263-book curriculum, because the remainder cannot be ingested at any rate.

Two-thirds of the Hindi curriculum is rendered unreadable by every PDF text extractor tested. NCERT’s older Hindi textbooks are typeset in **Walkman-Chanakya 905**, a legacy font mapping Devanagari glyphs onto ASCII codepoints with no `ToUnicode` table. Text extraction returns raw byte values. The analysis covers 41 of the 54 Devanagari-only books in the taught curriculum. The method samples a single chapter PDF per book, and for the remaining 13 books NCERT publishes no individual chapter files, returning HTTP 404 for chapters 1 through 3 at the documented URL scheme. Those 13 are absent for lack of a retrievable sample, not by selection, and no result here is conditioned on their contents (Table II):

TABLE II
HINDI CORPUS CORRUPTION ANALYSIS

Extraction Category	Books	Devanagari
Legacy font (Walkman-Chanakya)	27	0.0%
Unicode font (Kokila, Arya)	14	97.9–100%

V. CROSS-LINGUAL RETRIEVAL

Every retrieval and grounding result in this section and the next was measured against a single corpus state, recorded here so that the numbers are interpretable and reproducible. The database held **153 textbooks**, 138 English and 15 Hindi, comprising 1,542 chapters and **42,995 indexed passages**, with an embedding present for every passage (42,995 of 42,995; a passage lacking one is invisible to retrieval and is not otherwise detectable). All twelve classes are represented, ranging from 3 books at class 1 to 34 at class 12.

This is 153 of the 263-book taught curriculum. The remaining 110 are not pending but unobtainable. NCERT publishes no archive for most of them, one class 5 archive is corrupt, and the rest are the legacy-font Hindi readers of Section IV, which the pipeline rejects instead of storing as gibberish. Each is recorded with its reason, so the shortfall can be audited




Fig. 2. The five retrieval configurations over 22 queries, with 95% Wilson score intervals [17]: BM25 0/22, with rewriting 10/22, dense 11/22, dense with rewriting 14/22, and the deployed pooling of both 16/22. The intervals overlap across the three dense configurations, which is why this paper claims an ordering and not a margin. The lexical baseline’s interval excludes all three, and no sample-size argument softens that.

rather than asserted. An audit found two whose loss was this machine’s doing and not their source’s, and both have since been recovered, taking the corpus to 155 books. Every result here is measured against the 153-book state recorded above.

Composition is load-bearing for one result, not incidental to it. English science and mathematics books outnumber the Hindi readers roughly nine to one, which is what makes the grade-attribution result a test and not a formality. A retriever that defaults to the majority of its corpus would answer every Hindi-literature question from a science chapter.

Absolute cosine similarities should be read as characteristic of this corpus and not as a property of the method. Recall in particular is not estimated. There is no ground-truth relevant set for these books, and no such claim is made.

Retrieval was evaluated over **22 queries: two topics in each of the eleven supported Indian languages**, against the English-medium portion of the corpus. Holding the topic constant across languages makes language the only variable that moves. One topic uses ordinary vocabulary (the focusing of the human eye). The other is *photosynthesis*, whose name in most of these languages is a compound of morphemes meaning “light” and “joining”. The queries are committed with the benchmark, not only their results.

Five configurations were measured (Fig. 2), from a lexical baseline to the deployed system. A query counts as correct only when the *top* passage is on topic.

The lexical baseline retrieves nothing at all. BM25 [16] returns zero hits for all 22 queries, not merely poor ones. A Devanagari, Tamil or Perso-Arabic query shares no tokens with an English corpus, so the inverted index has nothing to match. Lexical retrieval is not weaker at this task, it is structurally incapable of it, and that is the argument for carrying a 568 M-parameter embedding model on a 4 GB machine. Given an English query from the rewrite it becomes viable at once (45%), which locates the difficulty precisely. The barrier is the script, not the vocabulary.

Multilingual RAG has been studied as a setting in its own right [9], including the tendency of such pipelines to answer in the language of the retrieved context instead of the question [10], which the separation of query language from answer language here is designed to prevent. The two




Fig. 3. Mean reciprocal rank of the deployed configuration for each language, averaged over the two topics. Rings mark 0.25, 0.5, 0.75 and 1.0; a language reaches the outer ring when both of its queries put an on-topic passage first. The two indentations that reach halfway are Urdu and Tamil, each of which answers one topic at rank 1 and returns nothing on photosynthesis.

arms fail on different inputs, and that is the result. Embedding the question directly handles ordinary vocabulary well (9 of 11) and collapses on compound terms (2 of 11), because the embedding attends to the constituent morphemes and retrieves optics for a word that means “light-joining”. Rewriting inverts both. It normalises the compound to its English canonical form and recovers 8 of 11, while losing three ordinary-vocabulary cases whose original phrasing had been more precise than its English paraphrase.

Neither is therefore the right configuration on its own. Pooling both retrievals scores **16 of 22**, above either arm, because the union recovers what each misses instead of averaging them. Fig. 3 gives the deployed configuration language by language, as the mean reciprocal rank of its two queries.

What 22 queries can and cannot establish. The intervals are wide and overlap for the three dense configurations. Pooling scores two queries above rewriting alone, at most four discordant pairs, which no test resolves at this n . We therefore claim only that pooling was never worse and that the ordering held across runs. One comparison is unambiguous. The lexical baseline’s $[0.00, 0.15]$ excludes every dense interval, and no sample-size argument rescues a method returning nothing on all 22 queries. Separating the narrower differences needs a query set in the hundreds, which takes native speakers of eleven languages, not compute.

The rewrite is a language-model call and is not deterministic at temperature zero. Repeated runs of this benchmark returned 16 and 17 of 22, so the pooled figure should be read as 16 ± 1 . The ordering of the five configurations was stable across runs, which is the claim the table is making. This is an empirical justification for a design that had been adopted on reasoning alone, and it also explains the earlier failure mode in which a mistranslated rewrite (“Rahim” rendered as “Rahman”) silently replaced a question that would have retrieved correctly.

Two queries fail in every configuration, both on photosynthesis. The Urdu term (*ziyai talif*, Perso-Arabic in origin) shares no morphology with either the English word or the Sanskrit-derived compounds the other nine languages use, so neither arm has a handle on it. The Tamil failure has a different cause:

the rewrite returned a paraphrase that dropped the term instead of normalising it, and direct embedding placed the passage fifth.

VI. MEMORY OPTIMIZATION

BGE-M3’s weights were evaluated at both full (float32) and half (float16) precision (Table III).

TABLE III
PRECISION COMPARISON FOR BGE-M3

Metric	float32	float16
Weight memory	2,166 MB	1,083 MB
Encode 4 queries	2,255 ms	1,197 ms
Model load time	8.5 s	8.1 s
Cosine fidelity (mean)	—	0.999998
Cosine fidelity (min)	—	0.999964

Half precision halves the weights and encodes $\sim 1.88\times$ faster. Fidelity, measured over 400 corpus passages rather than a handful of queries, is 0.999998 on average and 0.999964 at worst. Whether that matters is a question about rankings, so it was tested as one: over eight queries the top twelve passages are the same twelve under both precisions, in the same order for six of the eight, with the first difference at rank nine. Context is assembled from six passages, so nothing reported here is affected, but half precision can reorder a list rather than cannot.

Measurement against arithmetic. An earlier analysis summed component sizes to ~ 1.6 GB and concluded comfortable fit within 4 GB. A dedicated benchmarking script measures actual peak RSS through `getrusage (RUSAGE_SELF)`, a high-water mark surviving page eviction, unlike `ps` which reported 46 MB for a process holding a 1 GB model.

The actual peak is **2,094 MB**, not the 1,600 MB predicted. Loading model weights accounts for only 534 MB (`safetensors` memory-maps them). The first forward pass adds 1,560 MB as weight pages become resident and attention workspaces are allocated. The system fits within 4 GB with ~ 1.4 GB usable headroom (accounting for ~ 500 MB of operating system overhead), not the 2.4 GB arithmetic implied.

VII. GROUNDING AND HALLUCINATION DETECTION

A retrieval-augmented system can retrieve the correct chapter and still generate a fabricated answer, which is the failure mode surveyed across generation tasks by Ji et al. [20] and the one retrieval is supposed to remove. When asked about a specific story from the Hindi curriculum, the system retrieved the correct chapter, then described events that never occurred. The class was correct, the chapter was correct, the language was correct, and every detail was invented.

The **grounding score** computes cosine similarity between the generated answer and the source passages. Reference-free RAG evaluation frameworks score a comparable property, faithfulness, with a language model as judge [18]; the metric used here is a single cosine similarity, which costs one encoder pass and no second model. Lexical overlap cannot serve the




Fig. 4. Four grounding estimators over 16 matched and 240 mismatched pairs, with one standard deviation. What separates the populations is the distance between the two bars, not the height of either. Scoring against the single best passage raises both means without widening that distance, and scoring sentence by sentence narrows it.

purpose because the platform produces answers in a different language from the source passages. BGE-M3’s cross-lingual embedding property makes the metric functional in exactly this case.

Calibrating on a distribution. The threshold was originally set at 0.55 from four hand-verified answers, three correct at 0.583–0.701 and one fabricated at 0.489. Four points that happen to separate are not a calibration, and the difficulty is obtaining ungrounded answers in quantity. A model instructed to hallucinate does not fail the way a grounded pipeline fails.

The negatives are therefore constructed, not solicited. Every answer is generated normally from its own retrieved passages, then scored twice: against the passages it was written from (*matched*), and against another question’s passages (*mismatched*). A mismatched pair is precisely the failure the metric exists to catch (fluent, well-formed, on a different subject) and it requires no human labelling and no second generation pass. Sixteen questions spanning subjects and classes yield 16 matched and 240 mismatched pairs.

Grounding need not be measured against the concatenation of the retrieved passages, so four estimators were compared over identical generations (Fig. 4). The two runs of the experiment, with independently generated answers, agreed to within 0.007 on every entry.

The simplest estimator is also the best. Scoring the answer against all retrieved passages joined separates the two populations by 0.360, and no refinement improves on it. Scoring against the single best passage raises both means without widening the gap, and scoring sentence by sentence narrows it, because a correct answer contains connective and pedagogical sentences that appear in no passage, and averaging over them dilutes the signal the metric depends on.

The boundary was in the wrong place. The distribution shows that 0.55 separates the populations but sits too low. Thirty-five of the 240 mismatched pairs, one drifted answer in seven, score above it and pass as grounded. At 0.715 none do, while no matched answer falls below it. Balanced accuracy rises from 0.927 to 1.000, and the number of correct answers needlessly regenerated stays at zero. The deployed threshold was raised accordingly. The original four points were not wrong, merely too few to locate the boundary.

That 1.000 should be read for what it measures. The nega-

tives are answers written from *another question's* passages, so the task the score is perfect at is separating an answer from a different subject: topical drift, which is the failure we observed in deployment. It is not a measurement of hallucination detection in general, and an answer that stays on its retrieved topic while inventing a detail inside it would score as grounded here. The threshold is calibrated for the failure it was built to catch, and we claim no more than that.

The embedder is not judging its own geometry. BGE-M3 retrieved these passages and then scored the answers against them, so the separation could be a property of one embedding space and not of the text. Re-scoring the same generations settles it. Character 3–5 gram TF-IDF, not neural at all, separates the populations by 0.359, and a distilled multilingual MiniLM by 0.471, against BGE-M3's 0.360 (balanced accuracy 1.000, 0.996, 1.000). The gap is a property of the answers. Only the threshold is particular to the embedder deployed.

VIII. DISCUSSION

The measurement-driven methodology reveals a recurring pattern. Intuitive engineering assumptions are frequently contradicted by measurement, and the contradictions are more informative than the assumptions.

Availability. The platform, the ingestion pipeline and every benchmark reported here, including the 22 cross-lingual queries in their original scripts and the grounding calibration, are available at <https://github.com/rachitranka25/SHIKSHASETU>. The corpus itself is not redistributed. The pipeline reconstructs it from NCERT's own published archives.

Rights and ethics. NCERT publishes these textbooks for free public download. The system stores only locally derived text and vectors, redistributed with neither the code nor this paper, and a deployment beyond research use should confirm the applicable terms instead of inferring them from availability. No human subjects took part and no student data was collected. Every measurement here is taken from the system itself.

Limitations. (1) The grounding negatives are constructed by mismatching answers to other questions' passages, not collected from observed hallucinations. This captures topical drift, which is the failure we have seen, but not an answer that stays on topic and invents a detail within it. (2) Cross-lingual retrieval is evaluated on two topics per language. Broader topic coverage would tighten the estimate. (3) 27 legacy-font Hindi books remain unprocessable. (4) Memory measurements have run-to-run variance of up to 414 MB.

IX. CONCLUSION AND FUTURE WORK

This paper presented Shiksha Setu, a local-retrieval AI tutoring platform whose pipeline addresses the complete NCERT catalog of 559 textbooks and which delivers instruction in eleven Indian languages on consumer hardware with as little as 4 GB of RAM. The reported evaluation rests on 153 ingested textbooks and 42,995 indexed passages spanning all twelve classes. Every performance claim has been substantiated through kernel-level measurement, and engineering failures have been documented alongside their corrections.

The principal contributions are: (1) corpus construction with legacy-font corruption detection affecting 66% of Hindi textbooks, (2) cross-lingual retrieval measured over 22 queries in eleven languages, with an ablation showing that pooling the raw and rewritten queries retrieves correctly for 16 ± 1 of 22 against 11 and 14 for either alone, and that a lexical baseline retrieves none, (3) a grounding metric separating grounded from drifted answers by 0.360 in cosine similarity over 256 pairs, at a calibrated threshold that admits none of the 240 negatives where the previously shipped value admitted 35, and (4) half-precision inference halving memory at a mean cosine fidelity of 0.999998. Reading support for Brahmic scripts and a security audit of the deployment are reported in the extended version.

Future work targets a human-in-the-loop pilot with practising teachers, negatives drawn from observed hallucinations instead of constructions, a cross-encoder reranker for the compound-term failures, and a Walkman-Chanakya to Unicode transliterator to recover the legacy-font books.

REFERENCES

- [1] D. Kakwani *et al.*, "IndicNLP Suite: corpora, benchmarks and pre-trained models for Indian languages," in *Findings of EMNLP*, 2020.
- [2] P. Lewis *et al.*, "Retrieval-augmented generation for knowledge-intensive NLP tasks," in *Proc. NeurIPS*, 2020.
- [3] V. Karpukhin *et al.*, "Dense passage retrieval for open-domain question answering," in *Proc. EMNLP*, 2020, pp. 6769–6781.
- [4] X. Ma, Y. Gong, P. He, H. Zhao, and N. Duan, "Query rewriting for retrieval-augmented large language models," in *Proc. EMNLP*, 2023.
- [5] R. Nogueira and K. Cho, "Passage re-ranking with BERT," *arXiv preprint arXiv:1901.04085*, 2019.
- [6] J. Chen, S. Xiao, P. Zhang, K. Luo, D. Lian, and Z. Liu, "BGE M3-Embedding: multi-lingual, multi-granularity text embeddings via self-knowledge distillation," *arXiv preprint arXiv:2402.03216*, 2024.
- [7] A. Gala *et al.*, "IndicTrans2: accessible machine translation for all 22 scheduled Indian languages," *Trans. Mach. Learn. Res. (TMLR)*, 2023.
- [8] S. Khanuja *et al.*, "MuRIL: Multilingual representations for Indian languages," 2021.
- [9] Z. Xu *et al.*, "Multilingual retrieval-augmented generation for knowledge-intensive task," *arXiv preprint arXiv:2504.03616*, 2025.
- [10] Y. Chen *et al.*, "Language drift in multilingual RAG: characterization and decoding-time mitigation," *arXiv preprint arXiv:2511.09984*, 2025.
- [11] Y. A. Malkov and D. A. Yashunin, "Efficient and robust approximate nearest neighbor search using HNSW graphs," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 42, no. 4, pp. 824–836, Apr. 2020.
- [12] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-Networks," in *Proc. EMNLP*, 2019, pp. 3982–3992.
- [13] P. Micikevicius *et al.*, "Mixed precision training," in *Proc. ICLR*, 2018.
- [14] Khan Academy, "Khanmigo: an AI tutor for learners and assistant for teachers," 2023. [Online]. Available: khanmigo.ai
- [15] Ministry of Electronics and Information Technology, Government of India, "BHASHINI: National Language Translation Mission," 2022. [Online]. Available: bhashini.gov.in
- [16] S. Robertson and H. Zaragoza, "The probabilistic relevance framework: BM25 and beyond," *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, pp. 333–389, 2009.
- [17] E. B. Wilson, "Probable inference, the law of succession, and statistical inference," *J. Amer. Statist. Assoc.*, vol. 22, no. 158, pp. 209–212, 1927.
- [18] S. Es, J. James, L. Espinosa-Anke, and S. Schockaert, "RAGAs: Automated evaluation of retrieval augmented generation," in *Proc. EACL: System Demonstrations*, 2024, pp. 150–158.
- [19] Y. Gao *et al.*, "Retrieval-augmented generation for large language models: A survey," *arXiv preprint arXiv:2312.10997*, 2023.
- [20] Z. Ji *et al.*, "Survey of hallucination in natural language generation," *ACM Comput. Surv.*, vol. 55, no. 12, art. 248, 2023.
- [21] Department of School Education and Literacy, "UDISE+ 2024–25 report on unified district information system for education plus," Ministry of Education, Government of India, 2025.
- [22] Office of the Registrar General, India, "Census of India 2011, Table C-17: bilingualism and trilingualism," Min. Home Affairs, 2011.